Step1 Team Project Multivariate Analysis

Adrian White, Cesar Conejo, Xavier Bryant

12/6/2020

Introduction data set

We have selected the CRASH-2 data set provided by the Bank of Injury and Emergency Research Data from the UK. It describes the outcome of a randomized controlled trial and economic valuation of the effects of tranexamic acid on death, vascular occlusive events, and transfusion requirement in bleeding trauma patients. Tranexamic acid reduces bleeding in trauma patients undergoing surgery but is an expensive treatment option. The trial's objective was to assess the effects and cost-effectiveness of an early administration of this medication.

Participants of the study were adults with, or at risk of, significant bleeding within 8 hours of injury. Sample randomization was determined by the allocation of an eight-digit sequence randomly generated by a computer. Patients and staff were masked to the treatment allocation of the tranexamic acid.Patients and staff were masked to the treatment allocation of the tranexamic acid. However, the data set reflects some information in case of the decease of patients. As a result, we create a categorical variable called *death* with the final states 1 in case of passing and 0 if the patient survives.

We have adjusted the original data set to remove some variables that were not relevant to our investigation. We have removed variables regarding the exact surgical procedures administered to patients, various IDs, and details on the patient outcome. We removed the health outcome columns because of complications regarding missing data, where the boolean structure of the columns relating to specific outcomes, like stroke or pulmonary embolism, left a large number of cases with missing values. Instead, we added a boolean variable for a general outcome of survival to assess the efficacy of the procedure, rather than looking at particular health outcomes in post-surgery for living patients.

We will be using variables regarding the sex, age, and injury of the patient as well as certain biometrics, like blood pressure, respiratory and heart rates, details on surgical blood transfusion, and a boolean variable on the survival of the patient. Our selection provides us with a balance of continuous and categorical variables, many of which are boolean, with minimal complications due to missing data. In summary, the data set consists of $n = 9497$ observations, with 11 columns, which $p = 8$ are quantitative and 3 are qualitative.

Moreover, the normal ranges of the biometric measurements are also added, in order to have a point of comparison with the observations present in the data set and in this way determine if they are abnormal with respect to the normal metrics.

Summary variables in the data set

The variables in this dataset are the following:

1. sex: (Boolean) The sex of the patient (Male/Female)

2. age : (Numerical) Age of the patient(Years)

3. injurytime: (Numerical) Hours since injury (Hours)

4. injurytype: (Categorical) Type of injury {Blunt, Penetrating, Blunt and Penetrating}

5. sbp: (Numerical) Systolic Blood Pressure (mmHg). Normal range for adults at rest: less 120 mmHg.

6. rr: (Numerical) Respiratory Rate (breaths per minute). Normal range for adults at rest: 12 - 20 breaths per minute.

7. cc: (Numerical) Central Capillary Refill Time (seconds). Normal range for adults at rest. Less than 3 seconds.

8. hr: (Numerical) Heart Rate (beats per minute). Normal range for adults at rest: 60 - 100 bpm.

9. ndaysicu: (Numerical) Number of days in ICU (days)

10. ncell: (Numerical) Number of Units of Red Call Products Transfused.

11. death: (Boolean) Indicator if the patient survived after the procedure

A summary of the data type is the following:

variable	type_variable	sub_type_variable
sex	Qualitative	Nominal
age	Quantitative	Continuous
injurytime	Quantitative	Continuous
injurytype	Qualitative	Nominal
sbp	Quantitative	Continuous
rr	Quantitative	Continuous
cc	Quantitative	Continuous
hr	Quantitative	Continuous
ndaysicu	Quantitative	Discrete
ncell	Quantitative	Continuous
death	Qualitative	Nominal

A review of the structure of the data set is the following:

```
## 'data.frame':    9497 obs. of  11 variables:
##  $ sex       : Factor w/ 2 levels "male","female": 1 1 1 1 1 1 1 1 1 2 ...
##  $ age       : int   50 30 40 19 27 16 29 41 56 37 ...
##  $ injurytime: num   1 1 2 3 0.5 1 1 0.5 0.5 8 ...
##  $ injurytype: Factor w/ 3 levels "blunt","penetrating",..: 1 1 2 2 2 2 1 2 1 2 ...
##  $ sbp       : int   75 70 60 90 90 90 116 120 60 104 ...
##  $ rr        : int   28 26 20 30 26 28 15 15 9 23 ...
##  $ cc        : int   5 6 5 5 5 2 3 3 3 5 ...
##  $ hr        : int   120 130 120 90 96 118 118 70 100 92 ...
##  $ ndaysicu  : num   0 6 2 9 7 0 7 7 23 2 ...
##  $ ncell     : num   1 2 4 2 1 1 16 8 4 4 ...
##  $ death     : Factor w/ 2 levels "0","1": 2 1 2 2 1 1 1 1 1 1 ...
```

A summary of the values in the data set are:

```
##      sex            age           injurytime                      injurytype
##   male  :7906   Min.   :14.00   Min.   : 0.100   blunt                 :5211
##   female:1591   1st Qu.:24.00   1st Qu.: 1.000   penetrating           :2937
##                 Median :31.00   Median : 3.000   blunt and penetrating:1349
##                 Mean   :34.66   Mean   : 3.094
##                 3rd Qu.:43.00   3rd Qu.: 4.500
##                 Max.   :96.00   Max.   :48.000
##      sbp             rr              cc              hr
##   Min.   :  4.00   Min.   : 2.00   Min.   : 1.000   Min.   :  3.0
##   1st Qu.: 80.00   1st Qu.:20.00   1st Qu.: 2.000   1st Qu.: 96.0
##   Median : 90.00   Median :22.00   Median : 3.000   Median :110.0
##   Mean   : 93.13   Mean   :23.46   Mean   : 3.438   Mean   :108.1
##   3rd Qu.:104.00   3rd Qu.:28.00   3rd Qu.: 4.000   3rd Qu.:120.0
##   Max.   :225.00   Max.   :91.00   Max.   :20.000   Max.   :220.0
##    ndaysicu         ncell         death
##   Min.   : 0.000   Min.   : 0.000   0:7672
##   1st Qu.: 0.000   1st Qu.: 2.000   1:1825
##   Median : 1.000   Median : 3.000
##   Mean   : 4.137   Mean   : 3.912
##   3rd Qu.: 5.000   3rd Qu.: 5.000
##   Max.   :58.000   Max.   :60.000
```

Finally, the list of different values by column is the following:

Table 2: Count of distinct values of each variable

sex	age	injurytime	injurytype	sbp	rr	cc	hr	ndaysicu	ncell	death
2	81	78	3	153	58	16	154	47	47	2

Visual Analysis

Univariate Analysis

First, we will review the distribution of the variables involved in the data set.

In the case of *age*, the Figure 1 reflects how this variable appears to be largely weighted to the left, with lower ages featuring more frequently than those that are greater, possibly reflecting that younger people often take more risk and work higher at-risk occupations, raising their chance of experiencing trauma involving bleeding.


Figure 1: Distribution of age variable

Figure 2 below shows the distribution of the variable for injury time (*injurytime*) We can see how this variable is highly skewed to the left with almost all values falling below ten minutes since the injury was experienced. This is likely due to the fact, that in cases of serious injury, victims are brought to the hospital quite quickly as is the case for this data set on trauma. We apply a *log* transformation to this variable `log_injurytime = log(injurytime + 1)`, similarly to many other variables we have, because of the weight of the variables to the left of the distribution, in order to make a more normalized distribution. After the *log* transformation, we can see that two observations remain quite distinct and far to the right, with a value of four. They appear as potential outliers to the distribution, as they have almost 4 times the median value.

For *sbp* (systolic blood pressure), the distribution is a fairly centrally balanced distribution around 90 mmHg. Most observations around the mean and then a reasonably tight distribution of those who differ, making a distribution that is fairly Gaussian by nature. Furthermore, most people are fairly young in the sample and therefore would have rates, at a healthy level, that deviant less from the norm. The distribution is shown below in Figure 3.
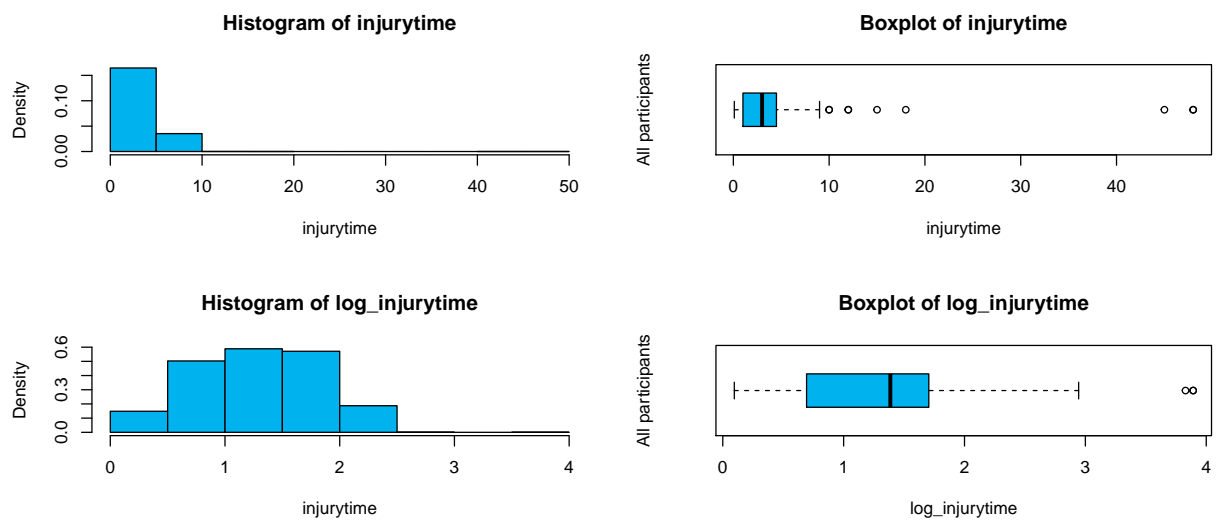
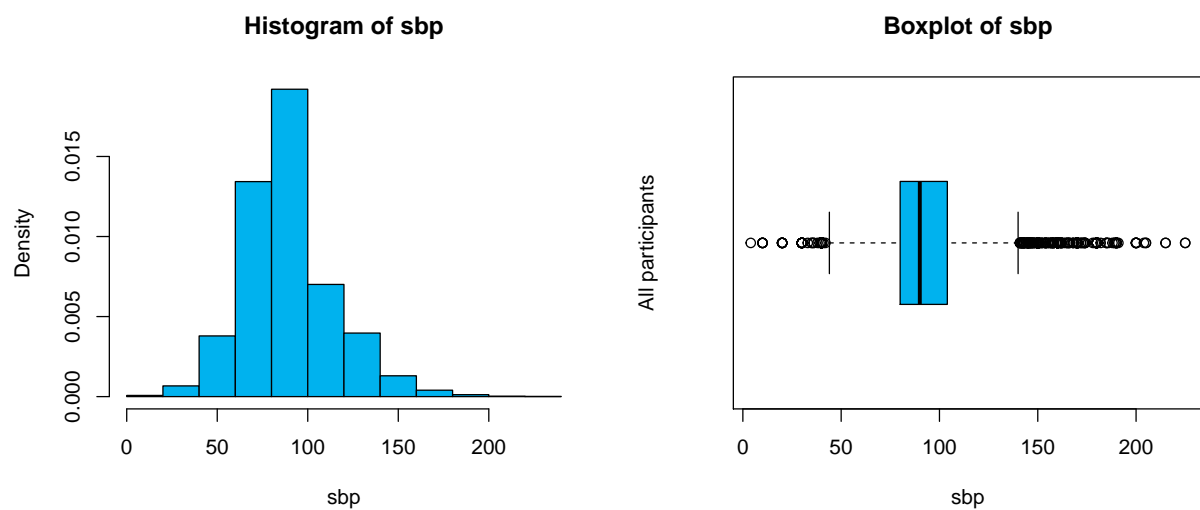Figure 2: Distribution of injurytime variable and log of injurytime

Figure 3: Distribution of sbp (Systolic Blood Pressure)

Rr (Respiratory Rate) appears, similar to sbp, to resemble a moderately balanced distribution around 22 respirations per minute, although *rr* is weighted more to the left. The distribution of this variable is shown in Figure 4. Taking a *log* transformation `log_rr = log(rr)`, we have the new distribution below as part of Figure 4. After the log transformation, the distribution remains quite spread out with a number of values beyond the whiskers of the box plot. The *log* transformation, in this case, has not significantly changed the shape of the distribution, it does not shift much more towards a normal distribution after the *log* transformation.
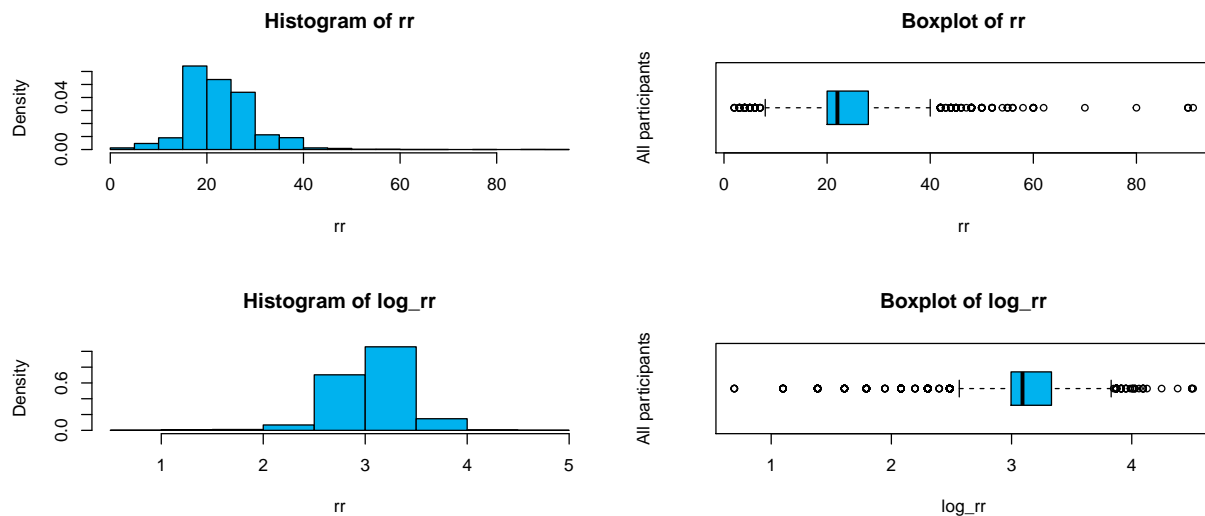

Figure 4: Distribution of rr (Respiratory Rate) and logrr

In the case of *hr* (Heart rate), Figure 5 has a distribution that seems fairly balanced at around 110, similar to the variables above, like *sbp* and *rr*. However, many values remain outside the whiskers and the IQR is quite tight, showing most people fall within a tight range but there is a number of people who have irregular rates on either side.


Figure 5: Distribution of hh (Heart Rate)

For *cc* (Central capillary refill time in seconds) has 75% of the observations below roughly 4 as we can appreciate in Figure 6. However, the distribution is right-skewed. As a result, we apply a *log* transformation `log_cc = log(cc)` that is given in the below part of Figure 6. The log has made the distribution more normal and was quite successful in this case. There remains a number of observations to the far right of the distribution, which could be possible outliers as they represent values approximately three times the median.
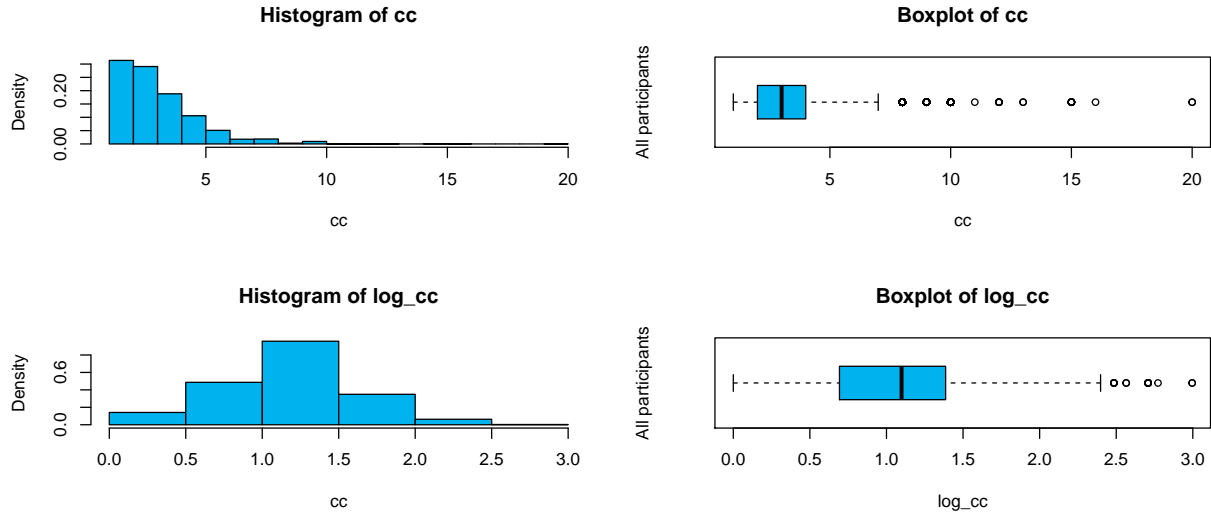

Figure 6: Distribution of cc (Central capillary) and log transformation

The following Figure 7 shows the distribution of the *ndaysicu:* variable. In this case, the distribution is heavily weighted to the left and right-skewed. Most patients it seems, with trauma injuries at high risk of bleeding, do not often need to remain in the hospital for long. The transformed distribution `log_ndaysicu = log(ndaysicu + 1)` is given in the below part of Figure 7. After the *log* transformation, this distribution doesn't appear much more Gaussian in nature. The values are still heavily weighted on the left. No values appear outside the whiskers, however, therefore this leads us to believe that there is not much likelihood of outliers in this distribution.


Figure 7: Distribution of ndaysicu and log transformation

7

Finally, the *ncell* distribution is weighted to the left with a median of 3 as we can see in Figure 8. The conclusion of this is that many patients only need a small number of or zero units of red cell products transfused. Due to this variable being highly weighted to the left, we apply the *log* transformation `log_ncell = log(ncell + 1)`. After the transformation, the variables visually appears to be more Gaussian, but a large number of values remain outside the whiskers on the right of the distribution. None, however, appear to be distinct enough from the others to be outliers.
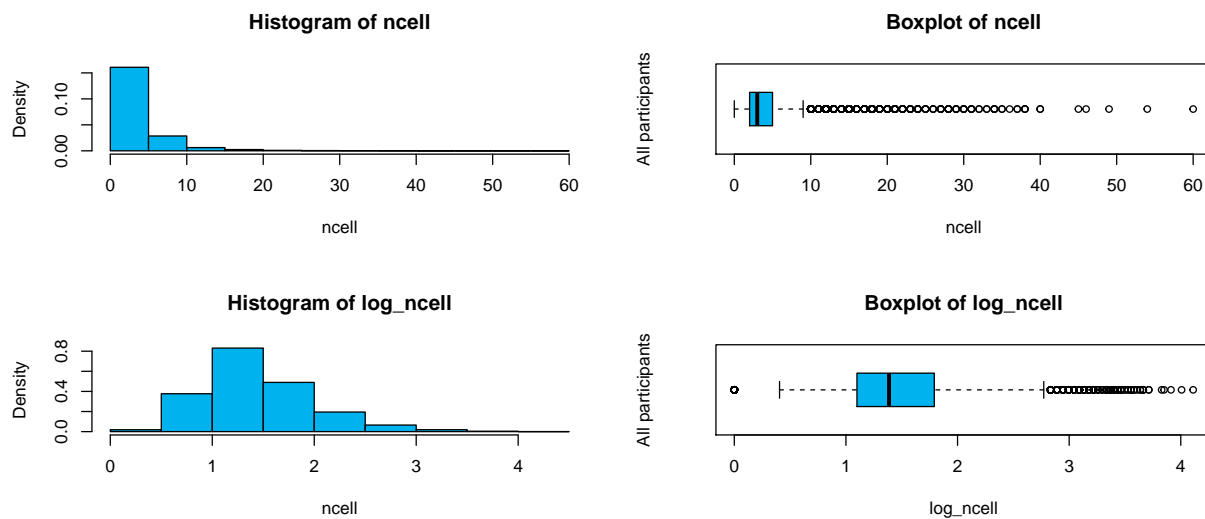

Figure 8: Distribution of ncell

Our primary categorical variable will be the distributions of *death*, where 0 represents people who survive and 1 for those who do not survive. Figure 9, shows that approximately for each death, 4 people survive. In the context of this problem, if we have an unbalanced proportion of people that survive, it can be considered as a sign that most people survive trauma injuries in the data set and that with the administration of the drug, most people survive. Although, this is not necessarily due to the treatment with the drug as both the control and treatment sample are contained in this data set, and were not separated (due to the absence of the treatment variable).

Figure 9: Distribution of deaths

Univariate Analysis by death - survival patients

Our next step is to study some relations of the quantitative variables in terms of our central categorical variable *death*.

Figure 10, shows us that those who survive (0) are on average younger than those who do not survive, which is logical, as younger people likely fair better in trauma accidents where there is a risk of significant bleeding (which this trail assesses). Aside from the median being slightly smaller than for those who survive, the distributions are fairly similar with a skew to the right. Participants in this trial are on average are quite young, roughly 34.7, which could be due to sampling, but also because younger people are more likely to suffer from trauma accidents due to having jobs with an increased propensity of injury and also more generally for risk-taking behavior.
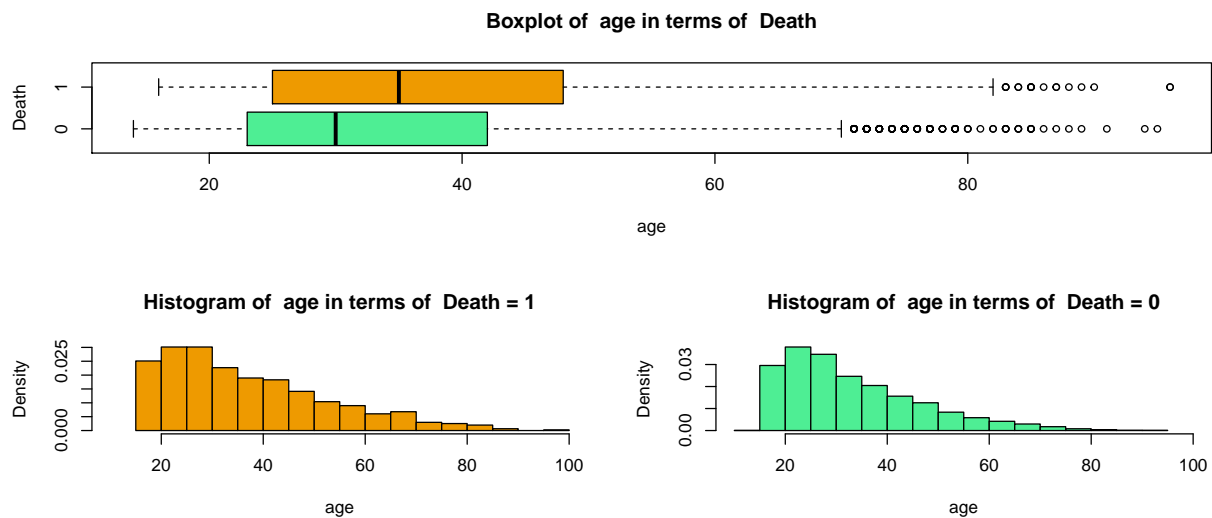

Figure 10: Distribution of Age in terms of death

Figure 11 showing *log injury time* demonstrates that those who survive (1) have longer relative injury time on average. The IQR for those who survive is also wider as well as the whiskers. However, there are still values of longer injury times for those who do survive, however they appear less frequently. This is slightly counterintuitive as we would expect those who took longer to get to the hospital or those who have been suffering longer, would be at a higher risk of *death*. However, it's possible that these people have less serious injuries, and have less risk of dying, and are then rushed to hospital with less urgency, having a long injury time.

Figure 11: Distribution of the log injurytime in terms of death

Figure 12 showing *sbp* (systolic blood pressure) against *death*, shows that those who survive (O), have distribution slightly to the right of that of the those who do not survive. This could be logical as those who survive have a higher blood pressure after the accident than those who eventually do not because their bodies are less weakened by the trauma.
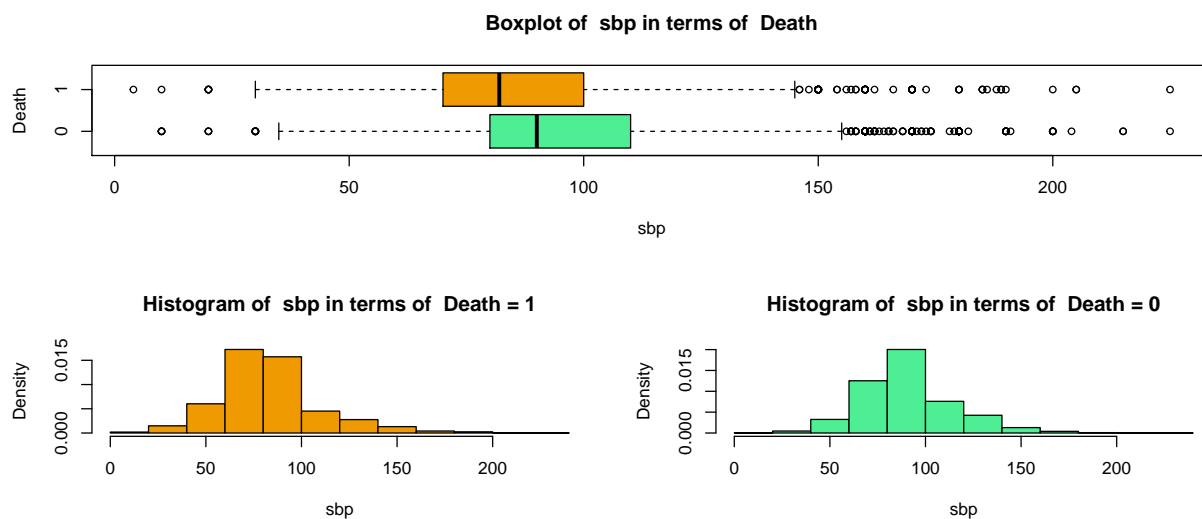
Figure 12: Distribution of sbp (Systolic Blood Pressure) in terms of death

Figure 13 compares *log respiratory rate* (rr) with those who do not survive (1) have a slightly higher mean than those who survive (1). This seems to lead us to believe, that those who do not survive, more often have a higher than average respiratory rate.

**Boxplot of log_rr in terms of Death**

Figure 13: Distribution of log transformation of rr (Respiratory Rate) in terms of death

Figure 14, similar to *log respiratory rate*, the median of the distribution of *log heart rate* for those who do not survive (1) is slightly higher than those who survive (0).
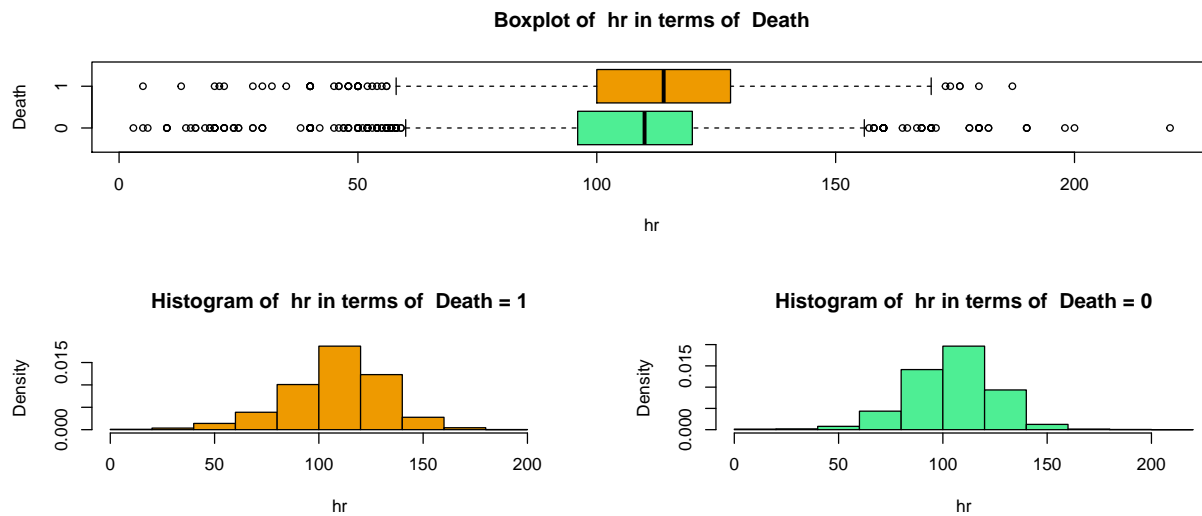
Boxplot of hr in terms of Death

Figure 14: Distribution of hh (hearth Rate) in terms of death

11

Figure 15, also resembling *log heart rate (hr)* and *log respiratory rate (rr)*, *log central capillary (cc)* (central capillary refill time in seconds) has a distribution further to the right, even more extreme than *hr* and *rr*, which could show that *cc* increases like *hr* and *rr* on average for people who will not survive. We also remain to see a limited number of observations to the extreme right, which could be outliers.


Figure 15: Distribution of log transformation of cc (Central capillary) in terms of death

Figure 16 shows that days in the ICU (*ndaysicu*), has a very similar distribution with an identical median, likely due to the fact that most people who suffer from trauma injuries only spend a very limited relative amount of time in the ICU. We see that, for those who survive (0), the 75% quartile is slightly larger, and when reviewing the histogram, more observations are on the further right end of the spectrum (more relative days in hospital). This possibly indicates to us that those who are more likely to survive are kept in the ICU for longer to heal in some cases.


Figure 16: Distribution of log ndaysicu in terms of death

Figure 17 shows the *log* value of red blood cells transfused (*ncell*) accounting for *death*. Those who do not survive (1), have a broader distribution with the 75% quartile being larger, possibly indicating that for those who will not survive, more blood cells are transfused for more grave injuries that are less survivable.
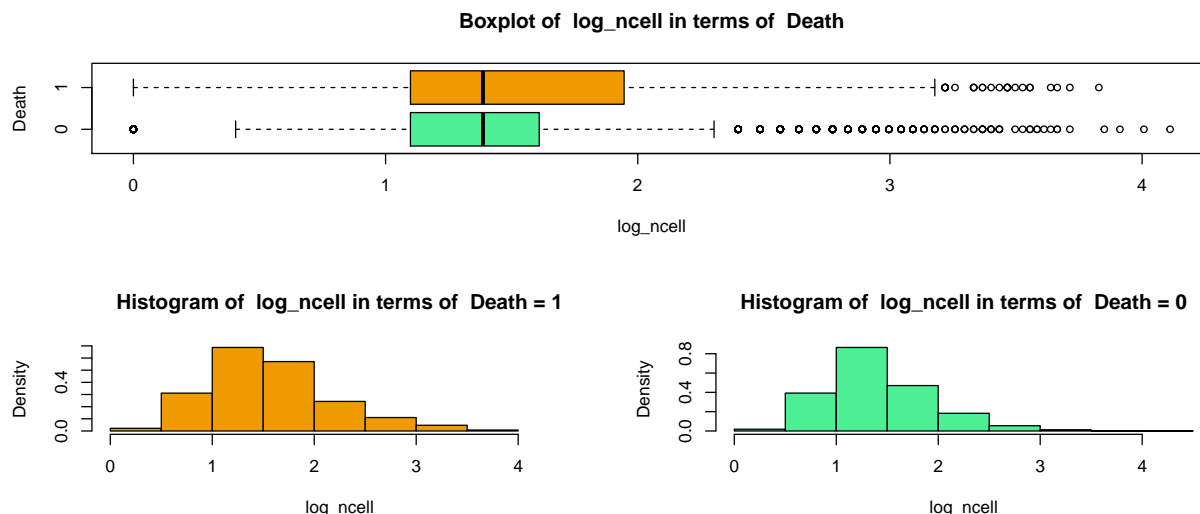

Figure 17: Distribution of log ncell in terms of death

Finally, in the scatter plot in Figure 18 we cannot see any significant relationships but the variables, while accounting for Death. There does not seem to be any visual relationship between any of the variables that are outstanding.

Furthermore, the PCP plot in Figure 19, generally matches with our observations about the scatter plots. Those who survive and do not, follow the same visual pattern. For *log_cc*, *hr*, and *log_ndaysicu*, those who survive appear to have values that are more broadly dispersed, but this could be due to the fact those who survive represent a majority of the sample size, making the visual depiction incorrect. Although comparing with the histograms and box plots, *log_cc*, *hr*, and *log_ndaysicu* does appear to have more dispersed distributions and wider IQRs.

Figure 18: Scatter plot of all quantitiative variables

Figure 19: PCP plot of all quantitiative variables

The Andrews' plot in Figure 20, matches our previous analysis in Figure 18 and 19, that variables share similar distributions and relationships when accounting for *death*, as the Andrew Plot does not show any dramatic differences between the two groups.

Andrews' Plot for individuals in Crash2 trial in terms of death

Figure 20: Andrews' plot of all quantitiative variables

In conclusion, we see that when accounting for different categorical variables and *death*, the relationships between our quantitative variables do not change significantly overall, looking at the scatter plots and other visuals. There doesn't appear to be any collinearity, whether *death* is accounted for or not. We can see some minor changes in distributions and means when dividing the distribution for those who do not and do survive, however, there is nothing too drastic. The IQRs for all our variables overlap, for example. The medians are even the same for *log_ncell* and *log_ndaysicu*. The variable with the largest difference appears to be, *log_cc* where the median for those who survive is at the 25% quartile of those who do not survive. For *sbp*, the median for those who do not survive is similar to the 25% quartile of those who survive.

In the next section, we will make inferences using sample estimators of mean, covariance, and correlation.

Sample Estimators

Sample mean

Below is the sample mean vector for all the variables in our analysis, without controlling for *death*.

Table 3: Sample mean

age	sbp	hr	log_injurytime	log_rr	log_cc	log_ndaysicu	log_ncell
34.66474	93.12909	108.0621	1.26759	3.106722	1.116876	0.9976935	1.394665

This table shows the mean while setting *death* to 1, or for those who do not survive.

Table 4: Sample mean: Do not survive

age	sbp	hr	log_injurytime	log_rr	log_cc	log_ndaysicu	log_ncell
37.82849	87.33808	111.4559	1.240578	3.110272	1.232432	0.9390286	1.524647

Now provided is the sample mean for the variables when the patient survives, or when $death = 0$.

Table 5: Sample mean: Survive

age	sbp	hr	log_injurytime	log_rr	log_cc	log_ndaysicu	log_ncell
33.91215	94.50665	107.2548	1.274015	3.105878	1.089387	1.011649	1.363745

Although touched upon above in our analysis of the box plots and histograms, comparing the means when controlling for *death*, overall, we do see some differences in the variables. The means of those who survive are more similar to the overall sample mean as those who survive represent more of the sample population. In terms of the difference between means we see that *age* is slightly higher for those who do not survive, which is logical, as those who are younger will likely be of better overall health and survive a trauma. *Hr*, *log_rr*, *log_cc* and *log_ncell* are all higher, and *sbp* lower, for those who do not survive which could be various responses related to the bodies response to more serious injury and it's medical treatment (more blood given, for example). *Log_injurytime* and *log_ndaysicu* are higher for those who survive, possibly as those who have less serious injuries have a higher likelihood to survive and take longer to get to a hospital. For similar reasons, *log_ndaysicu* could be longer for those who survive as those with more serious injuries were not able to survive for long, but those who could survive stayed longer in the ICU.

Sample covariance

Below are the sample covariance matrices for the general sample, and then for those who do not survive (1) and those who survive (0). The diagonal is the variances of the variables and the other values outside of the diagonals identify the covariances for the corresponding variables of the intersection of the respective row and column.

Table 6: Sample covariance matrix

	age	sbp	hr	log_injurytime	log_rr	log_cc	log_ndaysicu	log_ncell
age	201.8020767	5.6687061	-26.9775900	0.5843356	0.0754442	0.3511496	1.1379316	0.1173186
sbp	5.6687061	604.1019094	-117.2846636	1.8258493	-0.3695028	-2.8828312	1.6117323	-1.9082923
hr	-26.9775900	-117.2846636	459.8457610	-0.0753408	1.3004116	0.8685547	0.9063157	1.4739129
log_injurytime	0.5843356	1.8258493	-0.0753408	0.2877629	-0.0037522	0.0126630	0.0863252	0.0034858
log_rr	0.0754442	-0.3695028	1.3004116	-0.0037522	0.1092305	0.0119420	-0.0260511	-0.0025718
log_cc	0.3511496	-2.8828312	0.8685547	0.0126630	0.0119420	0.2448313	0.0636306	0.0344845
log_ndaysicu	1.1379316	1.6117323	0.9063157	0.0863252	-0.0260511	0.0636306	1.1508972	0.1949716
log_ncell	0.1173186	-1.9082923	1.4739129	0.0034858	-0.0025718	0.0344845	0.1949716	0.3404679

Second, here, we provide the sample covariance matrix for patients who did not survive ($death = 1$).

Table 7: Sample covariance matrix: Did not survive

	age	sbp	hr	log_injurytime	log_rr	log_cc	log_ndaysicu	log_ncell
age	261.7057666	17.272379	-51.3011548	0.8025150	0.0945272	0.3294187	1.1929405	-0.3849756
sbp	17.2723792	705.035309	-116.7715375	2.4101714	0.4360450	-3.2088292	4.1185430	-2.2891556
hr	-51.3011548	-116.771537	564.8052079	0.0045026	1.7090123	1.1424702	0.7975679	1.9352979
log_injurytime	0.8025150	2.410171	0.0045026	0.2799950	0.0020301	0.0093817	0.0999836	-0.0265201
log_rr	0.0945272	0.436045	1.7090123	0.0020301	0.1899808	0.0014615	-0.0137661	-0.0143161
log_cc	0.3294187	-3.208829	1.1424702	0.0093817	0.0014615	0.2228162	-0.0160445	0.0149818
log_ndaysicu	1.1929405	4.118543	0.7975679	0.0999836	-0.0137661	-0.0160445	0.8604265	0.1140010
log_ncell	-0.3849756	-2.289156	1.9352979	-0.0265201	-0.0143161	0.0149818	0.1140010	0.4261349

The following below is the sample covariance matrix for patients who survived (*death* = 0).

Table 8: Sample covariance matrix: Survive

	age	sbp	hr	log_injurytime	log_rr	log_cc	log_ndaysicu	log_ncell
age	184.6367863	8.3060201	-24.3595684	0.5577010	0.0676092	0.2486948	1.1796601	0.1156599
sbp	8.3060201	570.3044793	-111.6340107	1.6410806	-0.5550393	-2.6086143	0.9158253	-1.5962997
hr	-24.3595684	-111.6340107	431.5566186	-0.0673383	1.1998768	0.6880413	0.9909257	1.2344841
log_injurytime	0.5577010	1.6410806	-0.0673383	0.2894326	-0.0050994	0.0143642	0.0826221	0.0116551
log_rr	0.0676092	-0.5550393	1.1998768	-0.0050994	0.0900403	0.0143148	-0.0289142	0.0000845
log_cc	0.2486948	-2.6086143	0.6880413	0.0143642	0.0143148	0.2461653	0.0845804	0.0347028
log_ndaysicu	1.1796601	0.9158253	0.9909257	0.0826221	-0.0289142	0.0845804	1.2191014	0.2164958
log_ncell	0.1156599	-1.5962997	1.2344841	0.0116551	0.0000845	0.0347028	0.2164958	0.3151667

Comparing covariance matrices is a bit challenging as the magnitudes are not standardized, hence we reserve much of our analysis of the covariances between variables to the correlation matrices below, where the values are standardized between -1 and 1. We can analyze the difference in the individual variances of the variables when accounting for *death* here as this information is lost in the correlation matrices. To do this we use another additional measure of dispersion is called the *generalized variance*, which is defined as the determinant of the sample covariance. In the case of the three previous covariance table, the generalized variance is given by the following table. The generalized variance appears to be highest for the entire population, which is logical as it has the highest sample size and contains both groups for *death*. Surprisingly, the variance for those who do not survive (1) is significantly higher than for those who survive (0), despite having a drastically smaller sample size. This leads us to believe that data for those who do not survive is much more widely spread, possibly demonstrating that their indicators differ, many being related to health, are further from the median, or standard, healthy levels, indicating less likely to survive. In fact, we see a higher variance only in those who survive for the only two indicators that are not related to biological measurements in *log_ndaysicu* and *log_injurytime* (although very marginally in this case). As noted previously, *log_ndaysicu* could vary more for those who survive as those who do not survive may have more serious injuries where they are not able to make it to ICU treatment. It's possible that those who enter the ICU are more likely to survive after entering despite how long they spend there.

Table 9: Generalized variance

All population	Did not survive	Survive
114978.7	313091.3	74406.57

Sample correlation

Below are the sample correlation matrices for the general sample, and then for those who do not survive (1) and those who survive (0). The values outside of the diagonals identify the correlations for the corresponding variables in their respective row and column.

Table 10: Sample correlation matrix

	age	sbp	hr	log_injurytime	log_rr	log_cc	log_ndaysicu	log_ncell
age	1.0000000	0.0162355	-0.0885593	0.0766800	0.0160691	0.0499569	0.0746681	0.0141536
sbp	0.0162355	1.0000000	-0.2225256	0.1384817	-0.0454874	-0.2370449	0.0611251	-0.1330613
hr	-0.0885593	-0.2225256	1.0000000	-0.0065495	0.1834860	0.0818573	0.0393963	0.1177952
log_injurytime	0.0766800	0.1384817	-0.0065495	1.0000000	-0.0211642	0.0477075	0.1500037	0.0111366
log_rr	0.0160691	-0.0454874	0.1834860	-0.0211642	1.0000000	0.0730252	-0.0734743	-0.0133361
log_cc	0.0499569	-0.2370449	0.0818573	0.0477075	0.0730252	1.0000000	0.1198711	0.1194407
log_ndaysicu	0.0746681	0.0611251	0.0393963	0.1500037	-0.0734743	0.1198711	1.0000000	0.3114691
log_ncell	0.0141536	-0.1330613	0.1177952	0.0111366	-0.0133361	0.1194407	0.3114691	1.0000000

Figure 21, shown here is the correlation plot for the entire sample.

Figure 21: Sample Correlation Crash2 Dataset

Reviewing Figure 21 for the general sample covariance matrix we see that most values are fairly close to zero meaning most variables do not have a linear correlation. The maximum absolute value we see is the correlation between *log_ncell* and *log_ndaysicu* at 0.3113. Other correlation coefficients that exceed the absolute value of 0.1 are *log_ndaysicu* and *log_injurytime*, *sbp* and *log_injurytime*, *log_rr* and *hr* as well as *hr* and *sbp* at -0.2225. This is expected as our scatter plots did not visually reveal any linear relationships. According to Ratner (2009), correlation coefficients with values below 0.3 are considered weak or minimal relationships. All of our correlation coefficients are within this category. *Log_ncell* and *log_ndaysicu* are the only ones to reach the threshold of a "moderate relationship". The fact that *log_ncell* and *log_ndaysicu* have a limited moderate relationship could possibly be due to the fact those who stay in the hospital for longer need more blood and those who have serious injuries need more blood and need to stay in hospital more

often. We very small positive correlations between most variables but appear that *sbp* has minor negative correlations with most variables as well as *hr*, although these correlations remain minor.

Below is the sample correlation matrix for those who did not survive (*death*=1)

Table 11: Sample correlation matrix: Did not survive

	age	sbp	hr	log_injurytime	log_rr	log_cc	log_ndaysicu	log_ncell
age	1.0000000	0.0402106	-0.1334354	0.0937500	0.0134059	0.0431388	0.0794978	-0.0364547
sbp	0.0402106	1.0000000	-0.1850471	0.1715408	0.0376765	-0.2560164	0.1672173	-0.1320676
hr	-0.1334354	-0.1850471	1.0000000	0.0003580	0.1649837	0.1018409	0.0361794	0.1247455
log_injurytime	0.0937500	0.1715408	0.0003580	1.0000000	0.0088021	0.0375608	0.2037027	-0.0767761
log_rr	0.0134059	0.0376765	0.1649837	0.0088021	1.0000000	0.0071034	-0.0340485	-0.0503150
log_cc	0.0431388	-0.2560164	0.1018409	0.0375608	0.0071034	1.0000000	-0.0366435	0.0486202
log_ndaysicu	0.0794978	0.1672173	0.0361794	0.2037027	-0.0340485	-0.0366435	1.0000000	0.1882687
log_ncell	-0.0364547	-0.1320676	0.1247455	-0.0767761	-0.0503150	0.0486202	0.1882687	1.0000000

The following is the sample correlation matrix for those who survived (*death*=0)

Table 12: Sample correlation matrix survival patients

	age	sbp	hr	log_injurytime	log_rr	log_cc	log_ndaysicu	log_ncell
age	1.0000000	0.0255965	-0.0862962	0.0762902	0.0165817	0.0368888	0.0786281	0.0151619
sbp	0.0255965	1.0000000	-0.2250216	0.1277329	-0.0774554	-0.2201623	0.0347328	-0.1190669
hr	-0.0862962	-0.2250216	1.0000000	-0.0060252	0.1924861	0.0667548	0.0432018	0.1058514
log_injurytime	0.0762902	0.1277329	-0.0060252	1.0000000	-0.0315884	0.0538137	0.1390921	0.0385896
log_rr	0.0165817	-0.0774554	0.1924861	-0.0315884	1.0000000	0.0961511	-0.0872717	0.0005018
log_cc	0.0368888	-0.2201623	0.0667548	0.0538137	0.0961511	1.0000000	0.1543962	0.1245894
log_ndaysicu	0.0786281	0.0347328	0.0432018	0.1390921	-0.0872717	0.1543962	1.0000000	0.3492684
log_ncell	0.0151619	-0.1190669	0.1058514	0.0385896	0.0005018	0.1245894	0.3492684	1.0000000

Figure 23 is the correlation matrices while accounting for *death*, as well as, the Figure 22 which is a heatmap for the differences in the correlations.

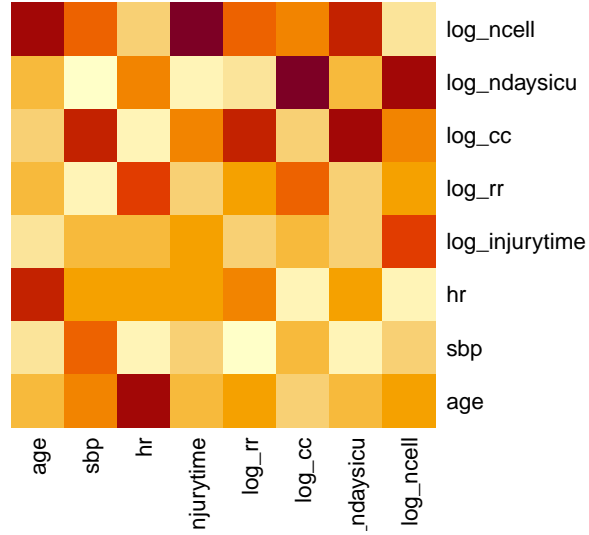**Heatmap: Difference of Matrix correlations death and survival**

Figure 22: Heatmap difference between correlation matrices

Figure 23: Sample Correlation for death and survival patients from Crash2 Dataset

We do not see any major changes in relationships for those who do not survive and people that do. Some variances do increase for those who survive such as *log_ndaysicu* and *log_ncell*, which almost doubles, and rises to a level of correlation at 0.3493 - considered moderate (Ratner, 2009). This may be because, for those that do not survive, they receive a lot of blood to attempt to save their lives, but do not survive long enough to stay in the ICU due to the severity of their trauma. This therefore would limit the covariance. For those who survive, they possibly could remain in hospital longer, therefore, receiving more blood over those days, increasing the covariance.

In conclusion, there is not much colinearity in these variables, while controlling for *death* or when not. This would mean, for example, that there is good potential for regression for the overall sample. However, though and more salient for our analysis at hand. the two categories for *death* are not easily separated, which means that drawing conclusions about the two separate groups, and comparing them, would be difficult and any statistical analysis would likely not be very accurate in drawing conclusions about the groups.

Outliers

Throughout our analysis, we did not have much reason to believe that outliers strongly affected our variables, we recognized some possibility in *cc* as well as *injurytime*. To compute the robust means we use the Minimum covariance determinant (MCD) estimators for the robust mean as we have many more observations ($n = 9497$) than quantitative variables ($p = 8$) (it also the reason why we did not perform shrinkage analysis). However, despite making many of the variable's distribution appear close to a Gaussian distribution our analysis with the MCD may be slightly flawed due to the fact that some of our variables are not Gaussian, even after transformation. Comparing the robust means to the regular we don't see any dramatic differences, confirming our analysis. For *cc* and *injurytime* we don't see much difference in fact, although these distributions are not Gaussian possibly diluting our ability to analyze. The largest differences are between *log_ndaysicu* and *log_ncell* although they are less than 0.05 of relative change (on the log-transformed variables). Overall, it appears that outliers are not very influential for any of the variables.

Table 13: Sample robust mean

age	sbp	hr	log_injurytime	log_rr	log_cc	log_ndaysicu	log_ncell
33.94433	92.32304	108.848	1.271321	3.137716	1.106643	0.9422688	1.357334

Table 14: Sample mean

age	sbp	hr	log_injurytime	log_rr	log_cc	log_ndaysicu	log_ncell
34.66474	93.12909	108.0621	1.26759	3.106722	1.116876	0.9976935	1.394665

Comparison of the eigenvalues of the covariance matrix for the sample and robust (MCD) means are given by Figure 24. This plot confirms our suspicions from above that those outliers do not affect the variables strongly as the eigenvalues for the two estimates do not differ strongly.

Analyzing the correlation plots for the sample and MCD estimator, we have further evidence that outliers are not important in our analysis, as the correlations hardly change. Some of the correlations appear stronger but this is only marginally and doesn't provide sufficient support to conclude that outliers would be influential in our analysis using any of our variables. We, therefore, end our examination of outliers here as the further investigation will likely not be helpful as outliers do not seem to be important.
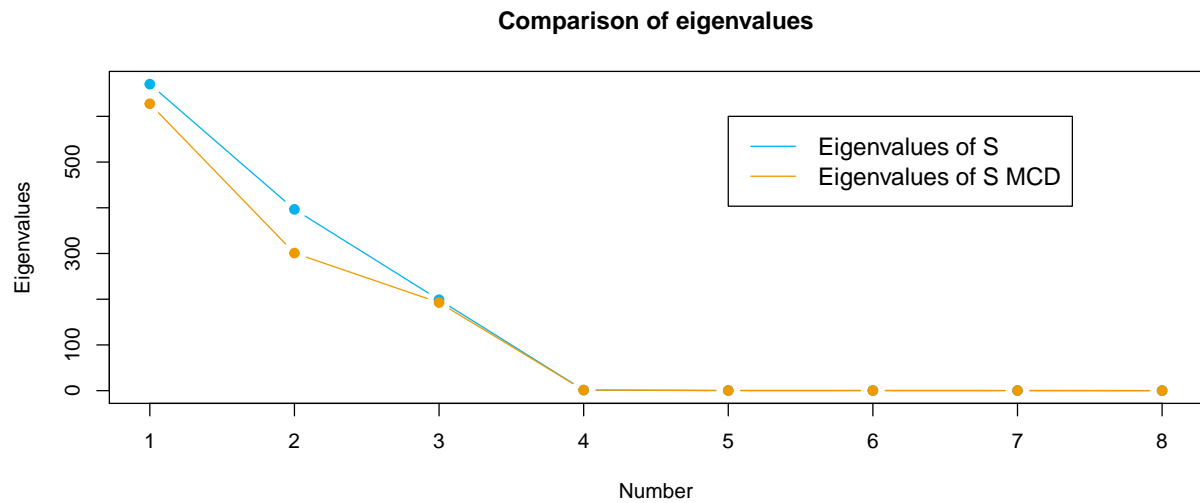
Figure 24: Comparison eigenvalues sample and robust covariance

Figure 25: Correlation all population and MCD correlation

Principal Component Analysis

The aim of Principal Component Analysis is to impose an alternative cartesian structure on our dataset that is optimized for dimensionality reduction. The data is restructured through an eigenvalue decomposition of the sample correlation matrix to derive the set of principal components. These feature vectors express the data in such a way that the explained variance of the data is concentrated in the first principal components. In other words, the descriptive capacity of our architecture is relegated to the top end of our sequence of dimensions. Organizing the data in such a way allows us to reduce the dimensionality of the set by removing the lower, less descriptive principal components while retaining a high level of variance.

In addition to providing the considerable benefits of dimensionality reduction, this approach allows us to consider underlying groups in the data. This will be our principal concern here since the number of dimensions in our case does not pose considerable complications and our ratio of columns to rows is very good. Our challenge, as previously discussed, is to determine patterns of behavior in our data.

In the previous sections, we have followed the necessary processes to prepare and normalize our data in order to optimize our PCA results. Non-standardized data may compromise PCA results, since the method is very sensitive to the variances of input variables. If the distributions differ considerably in range, the behavior of variables with greater variability will dominate our results. We have applied log transformations where necessary to account for skewed distributions. The R function prcomp() that we use to derive the principal components computes the mean-centered matrix and handles issues of scaling variables when necessary. By considering our sample correlations, we confirmed that we do not have issues of collinearity, which reduces the risk of redundant information in our variables.

Before the interpretation of the Principal Component Analysis, we can divide first our quantitative variables into three semantic categories:

1. Individual factors

   i) Age
   ii) injurytime

2. Biometrics

   i) sbp
   ii) rr
   iii) cc
   iv) hr

3. Medical attention

   i) ndaysicu
   ii) ncell

This is a useful exercise for us to express our own understanding of the descriptive nature of each of the variables. Each variable contributes a certain type of knowledge about the phenomenon we are studying and PCA reorganizes this information into general features. One drawback of PCA is that the principal components are less interpretable than our initial variables. Since the principal components are linear combinations of the initial variables, they don't have any real meaning and may only be understood as general directions of the data. As such, our analysis involves assigning possible interpretations of these directions.

PCA Complete dataset

In Figure 26, we plot the first two principal components considering the entire population. Immediately, we observe that there is no clear clustering of the data within the first two principal components. We have broken out the data by survival (green) and death (orange), but we see no clear linear relationship here either. Effectively, the full data set appears uniformly distributed about the origin. We remark that the density of orange points is more pronounced on the left quadrants, implying the proportion of death outcomes

PCA graph of individuals
First 2 PC by death

Figure 26: First two PCs based on the sample covariance all population

is somewhat greater for negative values of the first principal component, but uniformly distributed in the second.

We note that the first two principal components represent approximately 40% of the total variance explained in the model. This is less than half the explained variance of the model and may seem less than ideal, but as mentioned previously, our goal is to identify prominent clusters more so than building an optimal alternate model through dimensionality reduction. In Figure 30, we see that the captured variance for the sequence of eigenvalues does not exhibit a clear drop after the first few values. Instead, we see that the marginal explained variance decreases at a near constant rate. In fact, we see the greatest decrease in marginal explained variance after the second principal component. Thus, we will focus our analysis on the first two PCs.
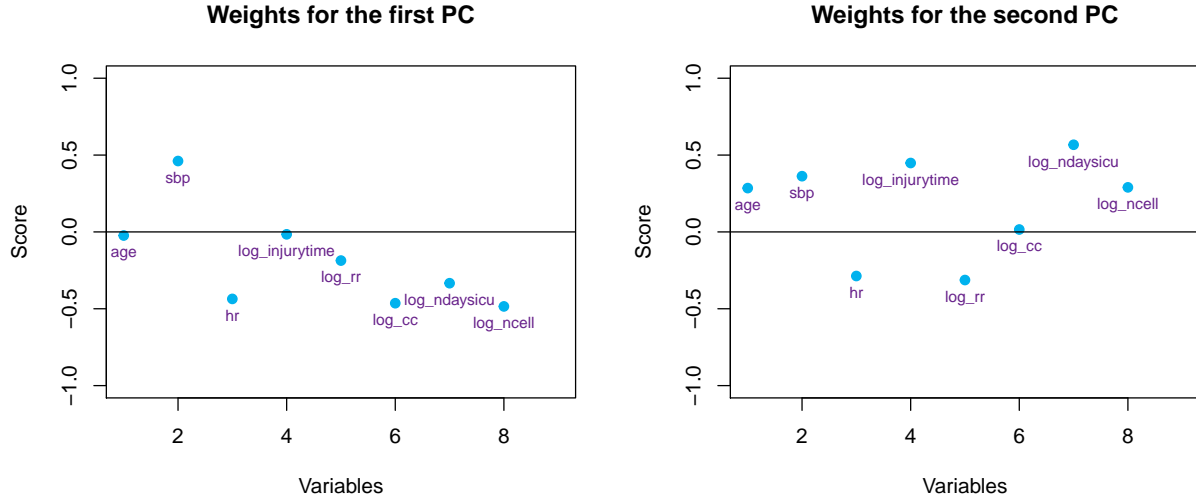
Figure 27: Loadings for the first and sencond PCs individually

Now, we consider the derivation of the first two principal components using loading plots, shown in Figure 27. These plots describe the weight of each variable on the respective principal component loadings. Weights with greater distances from zero imply a greater influence on the respective PC.

First Principal Component:

- We have variable sbp with a weight around +0.5 and hr, log_cc, and log_ncell with weights near -0.5, and log_rr, log_ndaysicu with smaller weights around -0.25.
- Variables age and log_injurytime have no weight on this PC.
- Our biometrics contribute the most weight to this PC, with sbp, hr, log_cc, and log_ncell with the most prominent weights. This feature may describe the effect of general patient health at the time of the accident.
- Some studies, such as Banegas et all and Kurl et all suggest that Systolic blood pressure is a more frequent cardiovascular risk factor than other blood-related measures and can be used to detect future diseases, which may explain its prominence.

Second Principal Component:

- We have log_ndaysicu and log_injurytime with weights near +0.5, and age, sbp, and log_ncell with weights around +0.25, and hr, log_rr with weights near -0.25.
- Variable log_cc has a weight of zero.
- We notice the contribution of age and log_injurytime to the second PC, while not present in the first.
- The variables log_ndaysicu and log_injurytime contribute the most weight to the second PC, so this feature describes the temporal context of the injury and recovery period.
- This may suggest that prompt medical attention increases chances for survival.

We continue our analysis of the PC loadings with Figure 28, which provides a 2D loading plot.


PCA graph of variables
Weights for the first two PCs

Figure 28: Loadings for the first two PC

We immediately note the most prominent variables: sbp, log_ndaysicu, log_ncell, and hr. Moreover, we may consider the geometry between the two PC axes. For this plot, the angles between the vectors provide information on how characteristics are correlated. Small angles imply positive correlations, large angles near 180° imply negative correlation, while orthogonality implies unlikely correlation.

We notice a possible negative correlation between sbp and the other biometrics log_rr, log_cc, and hr. In order to make conclusions, We consider possible explanations for this behavior based on medical studies:

1. Herakova et all report the relationship between the increase of blood pressure during exhalation and the decrease during inhalation. Moreover, the paper mentions that in general deep breathing could reduce blood pressure.

2. Harvard Heart Letter reports that an *isolated increase in blood pressure can drop the heart rate a little*. As a result, after a patient's trauma, the increase in *sbp* is followed by a decrease in heart rate, providing a possible explanation for the inverse relationship.

- We notice a strong positive correlation between *log_injurytime* and *age*, which we discussed in our previous sections. Greater age tends to be correlated with lower survival rates and longer treatment times.

- There is no relation between the *ndaysicu* with the biometrical measures as *sbp* and *hr*. It makes sense because this is a variable more related to medical attention. On the other hand *ndaysicu* is more related to the variables *ncell*, *injurytime*, and *age*. This implies a strong association between our individual characteristics and the medical attention prescribed.

In figure 29, we see a very similar visualization as in Figure 28, but with the observations included. Again, we notice a lack of clustering, but we have an understanding of the relationships of our characteristics.

Figure 29: PC Scores and PC loading all population

Figure 30 shows the explained variance of each eigenvalue. As previously discussed, the first two PCs represent approximately 40% of the explained variance, but failed to reveal any noticeable clusters. We see that the first four principal components explain approximately 63% of the model, this may provide more promising groupings.

We must note that the shape of this plot alerts us that our data may not be optimal to a standard PCA analysis. An ideal curve should be steep, with the majority of the variance concentrated in the first few PCs, then flatten out. Our curve, as discussed previously, decreases at a steady rate with a heavy right tail. Even the first four PCs only represent 63% of the variance, which is less than optimal. The method also failed to reveal any significant clusters. As a result, PCA may not be the most appropriate dimensionality reduction technique for our data.


Figure 30: Eigenvalues of the sample correlation matrix

The matrix scatterplot in figure 31 reflects the relations in the first four principal components. Here too, we do not appreciate a clear distinction between the groups (death and survival). Nor do we see linear relationships between the populations or a specific group. This provides further justification for our initial decision of considering only the first two PCs, since the subsequent two provide no further insight into possible clusters.

Finally, figure 32 explains the correlation between the variables and the principal components. This information is very similar to what we deduced graphically from the PC loadings plots. In the case of the first principal components, in magnitude, the biometrical measures and medical attention variables are more related to this component. In the case of the second principal component, the individual factors such as age and injury time acquire more relevance in the analysis.

We see that for the third and fourth PC's, *age* and *log_rr* represent the most prominent weights. In the case of PC4, the positive correlation between age and death becomes clear with the right skewed distribution of the orange points in the matrix scatterplot in figure 31.

Figure 31: PC Scores and PC loading all population

Figure 32: Correlation between dataset and all PC

PCA by Category

To further our analysis, we break out our data by survival outcome and develop the Principal component analysis for death (1825 individuals) and survival (7672 individuals) groups. In this case, we add the sex of the patients where females are the blue points and males are gray.

Similar to Figure 26, figure 33 shows there is no linear relation between the populations and the sex subgroups. Also, both principal components explain approximately 35% of the variance for each subgroup. This is similar to the PCA for all the population.

In the case of the loadings for the first principal component (figure 34) for the death population, the largest positive values are associated with *hr* and *cc*.

For the survival population, the first PC is greatly weighted by the variables associated with medical attention such as *ndaysicu* and *ncell*. As a result, we can observe the effect of receiving medical attention as a crucial factor for survival in case of accident or trauma. The variable sbp also has a significant negative weight.

The variable *injurytime* has no significant effect on the survival of the patients, but does weigh prominently in the first PC of the death group. We also notice that sbp has considerable weight for both groups.

When we compare this results with the result of the PCA for all the population, we appreciate that the effect of the medical attention is more relevant for the second principal component than for the first one.

Then, for the loading of the second principal component (figure 35), in the case of the death population, the effect of the medical attention is reflected with the fact that *ndaysicu* and *ncell* have the largest positive values. On the other hand, for the loadings of the second principal component for survival population, we can see how these variables are again relevant, but also are important the variables of *injurytime* and *sbp*. So, the second principal component for survival population reflects the effect of receiving faster medical help. This effect was also reflected in the second principal component for the general population.

Figure 33: PC Scores and PC loading all population

Figure 34: Loadings for the first PC by death and survival patients


Figure 35: Loadings for the second PC by death and survival patients

Figure 36 shows the PCA graph for loading and scores. For the Death patients, we notice the strong weight and positive correlation of *log_ndaysicu* and *log_injurytime*. A second prominent group of variables is *log_ncell*, *hr*, and *log_cc* which are strongly correlated together, but seemingly uncorrelated with the other group and negatively correlated with sbp. This aligns with our earlier observation, that elevated Systolic Blood Pressure may decrease heart rates.


Figure 36: Loadings for the first two PC by death and survival patients

For Survival patients, right panel of figure 36, we notice slightly different correlation patterns. For example, *log_ndaysicu* and *log_ncell* are positively correlated and significant for both PCs. For patients who survive, the duration and intensity of treatment are positively correlated. Meanwhile, for patients who do not survive, we see a correlation between the duration of treatment and the time of injury, independent of the intensity of treatment. For both groups, the intensity of treatment and time of injury are not correlated.

A variable that we would benefit from here is a measure of the severity of the trauma, which would provide a more nuanced understanding of these relationships. It is likely that a majority of those who survive suffered less intense wounds than those who died. This would explain why duration and intensity of treatment are correlated for survival patients, since this would most likely be indicative of progressively severe wounds. Meanwhile, for patients who do not survive, duration and intensity of treatment are not correlated, which may be because these patients generally suffered the most severe wounds and had low chances of survival. Thus, the speed at which they were able to begin treatment only bought them time, regardless of the treatment. If we had more information on the type and severity of the wounds, we would be able to test this hypothesis.

We have our PC Score plots with the observations mapped in figure 37. This is somewhat redundant and confirms our observations that there are no clear clusters and reiterates the correlations we discussed previously.

Finally, figures 38 and 39 provide us with the variance plots for the eigenvalues of the two groups. Similarly to the full population, we see that we have less than ideal reorganizations of the data using PCA on these two groups. The shape of the curves is almost linear, decreasing at a near-constant rate. This is not an ideal distribution of variance across principal components, where we would like a greater concentration at the top of the tail. We may want to consider alternative approaches to dimensionality reduction for this data.

Figure 37: PC Scores and PC loading all population for death and survival patiens


Figure 38: Eigenvalues of the sample correlation matrix for death patients

Figure 39: Eigenvalues of the sample correlation matrix for survival patients

Conclusions

The purpose of this study was to extract insights from our data using Principal Component Analysis. We selected a large data set and selected a subset of categorical and continuous variables for our analysis. Since we manually removed a number of variables, eliminating any pressing need for dimensionality reduction, our primary use for PCA was to identify any groups in the data and understand patterns regarding survival outcomes and treatment.

Our preliminary treatment of the data involved standardizing variables with significant skew using log transformations. We assessed the behavior of our continuous variables across different categories, finding strong patterns in age, injury time, and central capillary refill time in terms of survival outcomes. Using sample correlation matrices, we noticed recorded our strongest positive correlation between log_ncell and log_ndaysicu, relating the duration and intensity of treatments. We also noted negative correlations between Systolic Blood Pressure and other biometrics like Heart Rate and Respiratory Rate, which may be explained by the physiological effects of blood pressure on respiration.

Our application of Principal Component Analysis revealed no clusters in our data, even with breakdowns by sex and survival. However, we were able to more precisely assess the relationship between our variables and gain valuable insight into possible indicators of survival as well as multivariate patterns. We notice that for the general population, the first principal component is constructed mainly with the biometrics, while the second PC is weighted more by treatment times. This provides evidence for the importance of general biometric indicators, like blood pressure and respiratory rates, to understand the status of a patient.

We must note the sub-optimal distribution of variance across the principal components in the eigenvalue plots, which indicates that this method may not be appropriate for our data as is. Effectively, the behavior of our observations cannot be fully explained by a minimal number of feature vectors.

Through a categorical breakdown of PCA, we were able to assess differences in correlations between injury time, duration of hospitalization, and the intensity of treatment between patients who survive and those who do not. For survival patients, the duration and intensity of treatment are positively correlated, while patients who died show a positive correlation between the duration of treatment and the time of injury. In both groups, time of injury and intensity of treatment are not correlated.

This small divergence across groups may have alerted us to a broader and critical issue: existing patterns in our data are not adequately expressed by the variables we chose in our initial selection. It is our suspicion

that below our indicators for survival, there runs a current of behavior related to the *intensity of injury* that wields considerable influence on survival outcomes, the efficacy of treatment, and the biometrics. Unexpressed and dictatorial in its causality, this subliminal information may be the primary agent of sabotage in the restructuring of variance under the principal components, since the most profound agent of behavior remains hidden from us. This is easily grasped: the type and severity of treatment are fundamental in determining treatment options and are strongly correlated with the likelihood of survival.

Unfortunately, none of our selected variables provide even the most tangential information regarding the type and severity of the injury. Thus, this crucial information cannot even be inferred. However, the original data set did include such information. As such, we recommend that any further work on this analysis include a reconsideration of the selection of variables, where more information regarding the type and severity of injury be included in the data set.

References

1. Banegas JR, de la Cruz JJ, Rodriguez-Artalejo F, Graciani A, Guallar-Castillon P, Herruzo R. Systolic vs diastolic blood pressure: community burden and impact on blood pressure staging. J Hum Hypertens. 2002 Mar;16(3):163-7. doi: 10.1038/sj.jhh.1001310. PMID: 11896505.

2. S. Kurl, J.A. Laukkanen, R. Rauramaa, T.A. Lakka, J. Sivenius, and J.T. Salonen. Systolic Blood Pressure Response to Exercise Stress Test and Risk of Stroke. Sep 2001 https://doi.org/10.1161/hs0901.095395Stroke. 2001;32:2036–2041

3. Ratner, B. The correlation coefficient: Its values range between $+1/-1$, or do they? J Target Meas Anal Mark 17, $139-142$ (2009). https://doi.org/10.1057/jt.2009.5